Model Scoring & Governance Quick Reference

Interpreting scores and decisions from the deployed models · Knowledge base v1.0

1. Scores and decisions

A **model score** is a probability (0–1) produced by the deployed champion model. A **decision** is the business outcome after applying rules and thresholds to that score (approve / review / reject, route to queue, etc.). Always report which one you are quoting.

2. Interpretation bands

Score	Reading
0.90+	Very strong signal — the pattern closely matches the positive class the model was trained on.
0.72–0.89	Above the standard operational threshold (0.72) — actioned per the decision flow.
0.40–0.71	Grey zone — human review recommended; look at the record's strongest input factors.
Below 0.40	Weak signal — routine handling.

3. Governance rules

- Scores are decision support, never sole justification for adverse action against a person or supplier.
- Quote the model/decision name and version with every scored result.
- Out-of-range or missing inputs invalidate a score — the assistant must flag imputed or defaulted fields.
- Champion model changes are announced; historical scores are not comparable across champion versions.
- AutoML leaderboard metrics (AUC, KS, misclassification) are validation metrics — production performance is tracked separately.

4. When scoring fails

If a scoring call errors, the assistant reports the exact model input that was rejected and lists the expected fields. Re-ask with corrected inputs rather than approximating a score manually.